

## 6.SVD

April 24, 2025

### STAT 207: Singular Value Decomposition

Zhe Fei (zhe.fei@ucr.edu)

- NAS Chapter 9

In many modern applications involving large data sets, statisticians are confronted with a large  $m \times n$  matrix  $X = (x_{ij})$  that encodes  $n$  features on each of  $m$  objects.

- In gene microarray studies  $x_{ij}$  represents the expression level of the  $i$ th gene under the  $j$ th experimental condition.
- In information retrieval,  $x_{ij}$  represents the frequency of the  $j$ th word or term in the  $i$ th document.

The singular value decomposition (SVD) captures the structure of such matrices.

For a  $m \times m$  symmetric matrix  $A$ ,  $A = U\Sigma U^T$  with  $U = (u_1, \dots, u_m)$  gives

$$A = \sum_{j=1}^m \sigma_j u_j u_j^T.$$

When  $\sigma_j = 0$  for  $j > k$ ,  $A$  has rank  $k$ .

SVD generalizes the spectral theorem to nonsymmetric matrices.

$$A = \sum_{j=1}^k \sigma_j u_j v_j^T = U\Sigma V^T. \quad (1)$$

If  $A$  is  $m \times n$ , then write the SVD as

$$A = \begin{pmatrix} u_1 & \dots & u_k & u_{k+1} & \dots & u_m \end{pmatrix} \begin{pmatrix} \sigma_1 & 0 & \dots & 0 & 0 & \dots & 0 \\ 0 & \sigma_2 & \dots & 0 & & & \\ \vdots & \vdots & \ddots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & \sigma_k & 0 & \dots & 0 \\ 0 & 0 & \dots & 0 & 0 & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & 0 & 0 & \dots & 0 \end{pmatrix} \begin{pmatrix} v_1^T \\ \vdots \\ v_k^T \\ v_{k+1}^T \\ \vdots \\ v_n^T \end{pmatrix},$$

assuming  $k < \min\{m, n\}$ . The scalars  $\sigma_1, \dots, \sigma_k$  are said to be **singular values** and conventionally are listed in decreasing order. The vectors  $u_1, \dots, u_k$  are known as **left singular vectors** and the vectors  $v_1, \dots, v_k$  as **right singular vectors**.

## Basic Properties of the SVD

**NAS Proposition 9.2.1** Every  $m \times n$  matrix  $A$  has a singular value decomposition of the form (1) with positive diagonal entries for  $\Sigma$ .

Proof by induction.

Further we have

$$\begin{aligned} A^T &= \sum_{j=1}^k \sigma_j v_j u_j^T \\ AA^T &= \sum_{j=1}^k \sigma_j^2 u_j u_j^T \\ A^T A &= \sum_{j=1}^k \sigma_j^2 v_j v_j^T \end{aligned}$$

Hence,  $AA^T$  has nonzero eigenvalue  $\sigma_j^2$  with corresponding eigenvector  $u_j$ , and  $A^T A$  has nonzero eigenvalue  $\sigma_j^2$  with corresponding eigenvector  $v_j$ .

The following partial inverse is important in practice:

**NAS Proposition 9.2.2** The Moore-Penrose inverse  $A^- = \sum_{j=1}^k \sigma_j^{-1} v_j u_j^T$

enjoys the properties

$$(AA^-)^T = AA^-(A^-A)^T = A^-AAA^-A = AA^-AA^- = A^-.$$

If  $A$  is square and invertible, then  $A^- = A^{-1}$ . If  $A$  has full column rank, then  $A^- = (A^T A)^{-1} A^T$ .

**NAS Proposition 9.2.3** Suppose the matrix  $A$  has full SVD  $U\Sigma V^T$  with the diagonal entries  $\sigma_i$  of  $\Sigma$  appearing in decreasing order. The best rank- $k$  approximation of  $A$  in the Frobenius norm is

$$B = \sum_{j=1}^k \sigma_j u_j v_j^T.$$

Furthermore,  $\|A - B\|_F = \sqrt{\sum_{i>k} \sigma_i^2}$  and  $\|A - B\|_2 = \sigma_{k+1}$ .

## Applications

**Ridge Regression** In ridge regression, we minimize the penalized sum of squares

$$\begin{aligned} f(\lambda) &= \|y - X\beta\|_2^2 + \lambda \|\beta\|_2^2 \\ &= (y - X\beta)^T (y - X\beta) + \lambda \beta^T \beta. \end{aligned}$$

The gradient of  $f(\lambda)$  is

$$\nabla f(\lambda) = -2X^T(y - X\beta) + 2\lambda\beta.$$

Revised normal equations

$$(X^T X + \lambda I)\beta = X^T y,$$

with solution

$$\hat{\beta} = (X^T X + \lambda I)^{-1} X^T y.$$

If we further write  $X = \sum_j \sigma_j u_j v_j^T$ , then

$$X^T y = \sum_j \sigma_j u_j (u_j^T y), \quad X^T X + \lambda I = \sum_j (\sigma_j^2 + \lambda) v_j v_j^T.$$

The parameter estimates and predicted values reduce to

$$\hat{\beta} = \sum_j \frac{\sigma_j}{\sigma_j^2 + \lambda} u_j^T y v_j, \quad \hat{y} = X \hat{\beta} = \sum_j \frac{\sigma_j^2}{\sigma_j^2 + \lambda} (u_j^T y) u_j.$$

**Image Compression** An image (scene) is recorded as an  $m \times n$  matrix  $A = (a_{ij})$  of intensities.

- The entry  $a_{ij}$  represents the brightness of the pixel (picture element) in row  $i$  and column  $j$  of the scene.
- Storage issue when  $m$  and  $n$  are large
- Low rank approximate of  $B = (b_{ij})$

```
[1]: import numpy as np
from PIL import Image

def compress_image(image_path, k):
    # Load the image and convert to grayscale
    image = Image.open(image_path).convert('L')

    # Convert the image to a numpy array
    A = np.array(image)
    print('Original size', A.shape)

    # Apply the SVD to the image
    U, S, Vt = np.linalg.svd(A)

    # Truncate SVD matrices to retain only the k largest singular values
    U_k = U[:, :k]
    S_k = np.diag(S[:k])
    Vt_k = Vt[:k, :]

    # Reconstruct the compressed image
    B = U_k @ S_k @ Vt_k

    # Convert the numpy array back to an image
    compressed_image = Image.fromarray(B.astype('uint8'), 'L')

    return compressed_image
```

```
[2]: # Example usage
compressed_image = compress_image('cholesky.png', k=50)
compressed_image
```

```

-----
FileNotFoundError                                Traceback (most recent call last)
Cell In[2], line 2
      1 # Example usage
----> 2 compressed_image = compress_image('cholesky.png', k=50)
      3 compressed_image

Cell In[1], line 6, in compress_image(image_path, k)
      4 def compress_image(image_path, k):
      5     # Load the image and convert to grayscale
----> 6     image = Image.open(image_path).convert('L')
      7     # Convert the image to a numpy array
      8     A = np.array(image)

File ~/miniforge3/envs/stat_207/lib/python3.10/site-packages/PIL/Image.py:3227,
in open(fp, mode, formats)
    3224     filename = fp
    3226 if filename:
-> 3227     fp = builtins.open(filename, "rb")
    3228     exclusive_fp = True
    3230 try:

FileNotFoundError: [Errno 2] No such file or directory: 'cholesky.png'

```

```

[ ]: # Example usage
compressed_image = compress_image('cholesky.png', k=20)
compressed_image

```

Original size (920, 684)

```
[ ]:
```



**Principal Components** For a random vector  $Y$  with  $E(Y) = 0$  and variance matrix  $Var(Y)$ , the first principal component  $v_1^T Y$  is the linear combination that maximizes

$$Var(v^T Y) = v^T Var(Y) v.$$

With a centered random sample  $x_1, \dots, x_m$ , the sample variance is  $X^T X$  with

$$X = \frac{1}{\sqrt{m}} \begin{pmatrix} x_1^T \\ \cdot \\ \cdot \\ \cdot \\ x_m^T \end{pmatrix} = \sum_j \sigma_j u_j v_j^T.$$

The  $i$ th principal direction is given by the unit eigenvector  $v_i$ , and the variance of  $v_i^T x_j$  over  $j$  is given by  $\sigma_i^2$ .

### Jacobi's Algorithm for the SVD

By modifying the algorithm for eigen-decomposition, but without the need to calculate  $A^T A$ .

### Python Implementations

```
[ ]: import numpy as np

# generate a random matrix
A = np.random.randint(1, 10, size=(4,3))
A

[ ]: array([[6, 3, 5],
           [6, 8, 4],
           [3, 7, 4],
           [6, 9, 4]])

[ ]: # compute the SVD of A
U, s, Vt = np.linalg.svd(A)

print(s)
# check that U and Vt are orthogonal and s is a diagonal matrix
print(np.allclose(np.eye(4), np.dot(U.T, U)))
print(np.allclose(np.eye(3), np.dot(Vt, Vt.T)))

[19.36331759  3.96410117  1.53226432]
True
True

[ ]: # compute the eigenvalues and eigenvectors of A
B = A.T@A
w, v = np.linalg.eig(B)
```

```
print(w)
print(np.allclose(np.dot(w[0],v[:,0]), B@v[:,0]))
print(np.allclose(np.dot(w[1],v[:,1]), B@v[:,1]))
```

```
[374.93806795  2.34783396 15.71409809]
```

```
True
```

```
True
```

```
[ ]: w2, v2 = np.linalg.eig(A@A.T)
```

```
np.sum(w2)
```

```
[ ]: 392.99999999999955
```

```
[ ]:
```