

MCMC

June 6, 2023

1 STAT 207: Markov Chain Monte Carlo

Markov Chain Monte Carlo (MCMC) methods have a wide range of applications:

- Bayesian Inference: complex posterior distributions.
- Finance: option pricing, portfolio optimization, and risk assessment.
- Machine learning and AI: best moves.
- Physics and Chemistry: molecular dynamics simulations, protein folding, and quantum Monte Carlo methods.
- Genetics and Genomics, Social Sciences, and Many more.

Two major approaches:

- Hastings-Metropolis Algorithm
- Gibbs Sampling

1.1 Hastings-Metropolis Algorithm

Construct a Markov chain with a prescribed equilibrium distribution π on a given state space.

- Proposal stage: From stage i to stage j according to a probability density: $q_{ij} = q(j|i)$.
- Acceptance stage: $U \sim Unif[0, 1]$ compared to

$$a_{ij} = \min \left\{ \frac{\pi_j q_{ji}}{\pi_i q_{ij}}, 1 \right\}.$$

Remarks:

- If q is symmetric with $q_{ji} = q_{ij}$, then the acceptance probability reduces to

$$a_{ij} = \min \left\{ \frac{\pi_j}{\pi_i}, 1 \right\}.$$

- Check the detailed balance condition. Assume $\pi_i > 0$ for all i , wlog,

$$0 < \frac{\pi_j q_{ji}}{\pi_i q_{ij}} \leq 1$$

for some $i \neq j$. Then

$$\begin{aligned}\pi_i q_{ij} a_{ij} &= \pi_i q_{ij} \frac{\pi_j q_{ji}}{\pi_i q_{ij}} \\ &= \pi_j q_{ji} \\ &= \pi_j q_{ji} a_{ji}.\end{aligned}$$

- Irreducibility holds provided that $\pi_i > 0$ for all i and the proposal matrix $Q = (q_{ij})$ is irreducible.
- Aperiodicity: the acceptance-rejection step allows the chain to remain in place (Problem 4).

```
[1]: import numpy as np

def target_distribution(x):
    # Define the target distribution probabilities
    probabilities = [0.1, 0.2, 0.3, 0.4] # Example probabilities for discrete
    ↪distribution
    return probabilities[x]

def proposal_distribution(x):
    # Define the proposal distribution probabilities
    probabilities = np.full((4, 4), 0.25) # Example probabilities for discrete
    ↪distribution
    return probabilities[x,]

def hastings_metropolis(target_dist, proposal_dist, num_samples):
    # Initialize the current sample
    current_sample = 0
    samples = []

    for _ in range(num_samples):
        # Generate a proposal sample from the proposal distribution
        proposal_sample = np.random.choice([0,1,2,3],
    ↪p=proposal_dist(current_sample))

        # Calculate the acceptance ratio
        acceptance_ratio = min(1, target_dist(proposal_sample) /
    ↪target_dist(current_sample))

        # Accept or reject the proposal sample based on the acceptance ratio
        if np.random.uniform(0, 1) < acceptance_ratio:
            current_sample = proposal_sample

        # Save the current sample
        samples.append(current_sample)

    return samples
```

```

# Set the parameters
num_samples = 10000

# Run the algorithm
samples = hastings_metropolis(target_distribution, proposal_distribution,
    ↪ num_samples)

```

```

[2]: unique, counts = np.unique(samples, return_counts=True)
     print(np.asarray((unique, counts)).T)

```

```

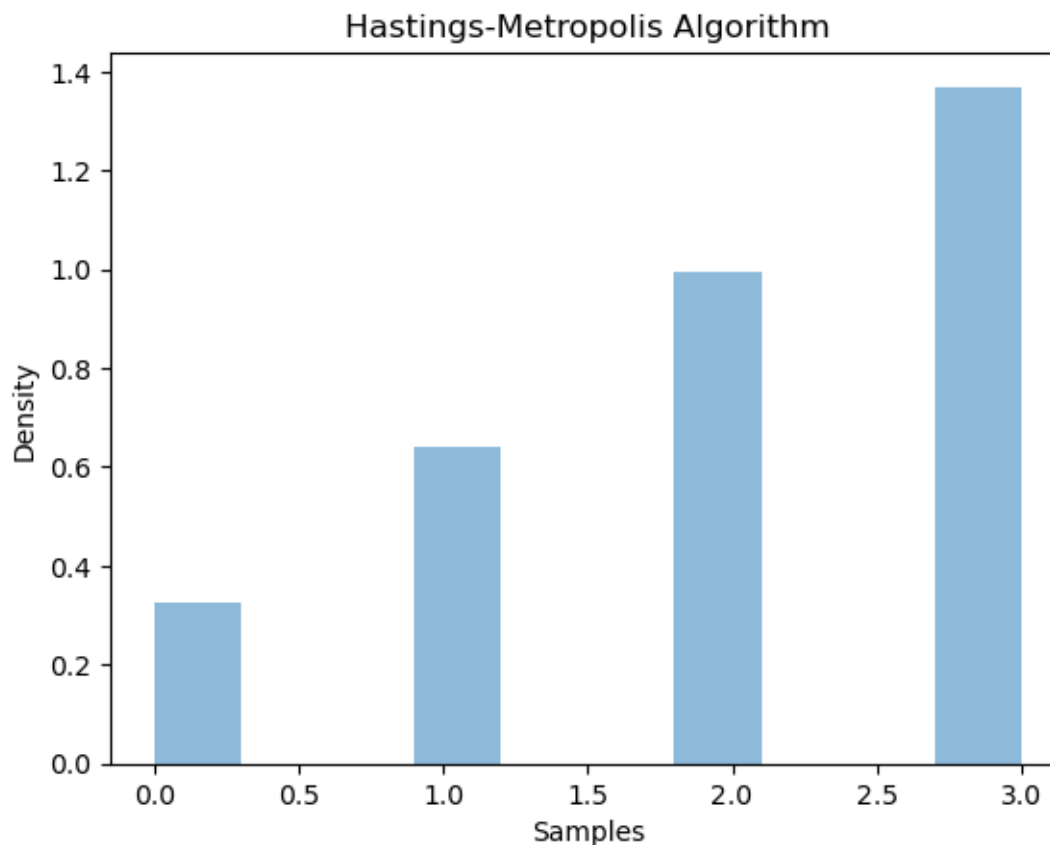
[[ 0  980]
 [ 1 1924]
 [ 2 2988]
 [ 3 4108]]

```

```

[3]: # Plot the samples
     import matplotlib.pyplot as plt
     plt.hist(samples, density=True, alpha=0.5)
     plt.xlabel('Samples')
     plt.ylabel('Density')
     plt.title('Hastings-Metropolis Algorithm')
     plt.show()

```



1.2 Gibbs Sampling

- The Gibbs sampler is a special case of the Hastings-Metropolis algorithm for Cartesian product state spaces.
- Our random variable of interest has d components, $x = (x_1, x_2, \dots, x_d)$,
- Suppose we can simulate the distribution of each component conditional on the others, i.e., $\pi(x_k | x_1, \dots, x_{k-1}, x_{k+1}, \dots, x_d)$ for all k .
- We want to sample from the joint distribution, $\pi(x)$.
- Gibbs sampling constructs a Markov Chain, $x^{(1)} \rightarrow x^{(2)} \rightarrow \dots$, with step $x^{(j)} \rightarrow x^{(j+1)}$ given by:

Simulate

- $x_1^{(j+1)}$ from $\pi(x_1^{(j+1)} | x_2^{(j)}, \dots, x_d^{(j)})$
- $x_2^{(j+1)}$ from $\pi(x_2^{(j+1)} | x_1^{(j+1)}, x_3^{(j)}, \dots, x_d^{(j)})$
- $x_3^{(j+1)}$ from $\pi(x_3^{(j+1)} | x_1^{(j+1)}, x_2^{(j+1)}, x_4^{(j)}, \dots, x_d^{(j)})$
- ...

- By iteratively updating each component, the Gibbs sampler generates a sequence of samples

that approximates the joint distribution of interest.

- To see this, for a component x_c transit from i to j , the transition probability is

$$q_{ij} = \frac{1}{d} \frac{\pi_j}{\sum_{j \in \{k: k_{-c} = i_{-c}\}} \pi_k},$$

which satisfies $\pi_i q_{ij} = \pi_j q_{ji}$.

- Denote $P^{(c)}$ the transition matrix for changing component c while leaving other components unaltered.
 - Random sampling: $R = \frac{1}{d} \sum_c P^{(c)}$;
 - Sequential sampling: $S = \prod_c P^{(c)}$.
 - $\pi R = \pi$ and $\pi S = \pi$.
 - R satisfies detailed balance while S ordinarily does not (Problem 6).

1.2.1 Conjugate distributions

A likelihood $p(x|\theta)$ and a prior density $p(\theta)$ are said to be conjugate provided the posterior density $p(\theta|x)$ has the same functional form as the prior density.

TABLE 26.1. Conjugate Pairs

Likelihood	Density	Prior	Density
Binomial	$\binom{n}{x} p^x (1-p)^{n-x}$	Beta	$\frac{1}{B(\alpha, \beta)} p^{\alpha-1} (1-p)^{\beta-1}$
Poisson	$\frac{\lambda^x}{x!} e^{-\lambda}$	Gamma	$\frac{\beta^\alpha \lambda^{\alpha-1}}{\Gamma(\alpha)} e^{-\beta \lambda}$
Geometric	$(1-p)^x p$	Beta	$\frac{1}{B(\alpha, \beta)} p^{\alpha-1} (1-p)^{\beta-1}$
Multinomial	$\binom{n}{x_1 \dots x_k} \prod_{i=1}^k p_i^{x_i}$	Dirichlet	$\frac{\Gamma(\sum_{i=1}^k \alpha_i)}{\prod_{i=1}^k \Gamma(\alpha_i)} \prod_{i=1}^k p_i^{\alpha_i-1}$
Normal	$\sqrt{\frac{\tau}{2\pi}} e^{-\tau(x-\mu)^2/2}$	Normal	$\sqrt{\frac{\omega}{2\pi}} e^{-\omega(\mu-\theta)^2/2}$
Normal	$\sqrt{\frac{\tau}{2\pi}} e^{-\tau(x-\mu)^2/2}$	Gamma	$\frac{\beta^\alpha \tau^{\alpha-1}}{\Gamma(\alpha)} e^{-\beta \tau}$
Exponential	$\lambda e^{-\lambda x}$	Gamma	$\frac{\beta^\alpha \lambda^{\alpha-1}}{\Gamma(\alpha)} e^{-\beta \lambda}$

1.2.2 Example: Ising Model

- Consider m elementary particles equally spaced around the boundary of the unit circle.
- Each particle c can be in one of two magnetic states—spin up with $i_c = 1$ or spin down with $i_c = -1$.
- The Gibbs distribution

$$\pi_i = \exp(\beta \sum_c i_c i_{c+1})$$

takes into account nearest-neighbor interactions in the sense that states like $(1, 1, 1, \dots, 1, 1, 1)$ are favored and states like $(1, -1, 1, \dots, 1, -1, 1)$ are less likely for $\beta > 0$.

- If we elect to resample component c , then the choices $j_c = -i_c$ and $j_c = i_c$ are made with respective probabilities

$$\frac{e^{\beta(-i_{c-1}i_c - i_c i_{c+1})}}{e^{\beta(i_{c-1}i_c + i_c i_{c+1})} + e^{\beta(-i_{c-1}i_c - i_c i_{c+1})}} = \frac{1}{e^{2\beta(i_{c-1}i_c + i_c i_{c+1})} + 1}$$

$$\frac{e^{\beta(i_{c-1}i_c + i_c i_{c+1})}}{e^{\beta(i_{c-1}i_c + i_c i_{c+1})} + e^{\beta(-i_{c-1}i_c - i_c i_{c+1})}} = \frac{1}{1 + e^{-2\beta(i_{c-1}i_c + i_c i_{c+1})}}$$

1.2.3 Example: Capture-Recapture Estimation

- To estimate the number of fish f in a lake.
- Fish t times, mark each fish caught. Data (c_i, r_i) .
- $u_i = \sum_{j=1}^r (c_i - r_i)$ unique fish.
- Independent binomial sampling with success probability p_i for trial i ,

$$\begin{aligned} & \prod_{i=1}^t \binom{u_{i-1}}{r_i} p_i^{r_i} (1-p_i)^{u_{i-1}-r_i} \binom{f-u_{i-1}}{c_i-r_i} p_i^{c_i-r_i} (1-p_i)^{f-u_{i-1}-c_i+r_i} \\ &= \binom{u_{i-1}}{r_i} \binom{f-u_{i-1}}{c_i-r_i} p_i^{c_i} (1-p_i)^{f-c_i} \\ &= \frac{f!}{(f-u_t)!} \prod_{i=1}^t \binom{u_{i-1}}{r_i} p_i^{c_i} (1-p_i)^{f-c_i} \end{aligned}$$

where $u_0 = r_1 = 0$.

- Poisson prior for f and independent beta priors on p_i 's. The joint density is

$$\begin{aligned} & \frac{f!}{(f-u_t)!} \prod_{i=1}^t \binom{u_{i-1}}{r_i} p_i^{c_i} (1-p_i)^{f-c_i} \frac{\lambda^f e^{-\lambda}}{f! B(\alpha, \beta)^t} \prod_{i=1}^t p_i^{\alpha-1} (1-p_i)^{\beta-1} \\ &= \frac{\lambda^f e^{-\lambda}}{(f-u_t)! B(\alpha, \beta)^t} \prod_{i=1}^t \binom{u_{i-1}}{r_i} p_i^{\alpha+c_i-1} (1-p_i)^{\beta+f-c_i-1}. \end{aligned}$$

- The Gibbs update for p_i is $Beta(\alpha + c_i, \beta + f - c_i)$.
- The Gibbs update for f is that $f - u_t$ following Poisson with mean $\lambda \prod_{i=1}^t (1-p_i)$.

1.3 Example: Slice Sampling

- X with density $f(x)$
- Gibbs sampling on the pair (X, Y) with auxiliary Y .
- Sample uniformly from the region under the graph of $f(x)$.
- Given X , Y is uniform from $[0, f(X)]$; given $Y = y$, X is uniform from $\{x : f(x) \geq y\}$.

$$\Pr(X \in A) = \int_A \frac{1}{f(x)} \int_0^{f(x)} dy f(x) dx = \int_A \int_0^{f(x)} dy dx.$$

- For example, $f(x) = e^{-x^2/2}$,
- When $f(x)$ is concave or log-concave, a top set is convex and therefore connected. Otherwise, it can be disconnected.

1.4 Example: Independence Sampler

- If $q_{ij} = q_j$, then candidate points are drawn independently of the current point.
- Want q_i close to π_i , introducing the importance ratios $w_i = \pi_i/q_i$ and the acceptance probability:

$$a_{ij} = \min \left\{ \frac{w_j}{w_i}, 1 \right\}.$$

1.5 Example: Random Walk

- If $q_{ij} = q_{j-i}$ for some density q_k , and $q_k = q_{-k}$, then the acceptance probability is

$$a_{ij} = \min \left\{ \frac{\pi_j}{\pi_i}, 1 \right\}.$$

- Consider the density on \mathbb{R}^2 :

$$f(x) = e^{-\frac{(\|x\|_2 - 1)^2}{2\sigma^2}} e^{-\frac{(x_2 - 1)^2}{2\delta^2}} = f_1(x)f_2(x)$$

- The proposed symmetric density:

$$g(y) = \frac{1}{2\pi\gamma^2} e^{-\frac{\|y\|_2^2}{2\gamma^2}}$$

```
[4]: import numpy as np
import matplotlib.pyplot as plt

def f(x, sigma, delta):
    norm_term = np.exp(-((np.linalg.norm(x, axis=0)-1)**2)/(2*sigma**2))
    exp_term = np.exp(-((x[1]-1)**2)/(2*delta**2))
    return norm_term * exp_term


# Define the range of x and y values
x = np.linspace(-1.5, 1.5, 100)
y = np.linspace(-1.5, 1.5, 100)

# Create a grid of (x, y) points
X, Y = np.meshgrid(x, y)

# Evaluate the function at each point of the grid
```

```
Z = f(np.array([X, Y]), sigma=0.1, delta=1)
```

```
# Plot the contours
plt.contour(X, Y, Z, levels=20)
plt.xlabel('x')
plt.ylabel('y')
plt.title('Contours of f(x)')
plt.colorbar()
plt.show()
```



```
[5]: def g(y, gamma):
      norm_term = np.exp(-(np.linalg.norm(y)**2)/(2*gamma**2))
      return norm_term / (2*np.pi*gamma**2)

def hastings_metropolis_sampling(target_density, proposed_density,
    ↪initial_sample, num_samples, gamma):
    samples = [initial_sample]
    accepted_samples = 0
```



```

    for _ in range(num_samples):
        current_sample = samples[-1]
        proposed_sample = np.random.normal(current_sample, gamma)

        acceptance_prob = min(1, (target_density(proposed_sample) *
        ↪proposed_density(current_sample, gamma)) /
                               (target_density(current_sample) *
        ↪proposed_density(proposed_sample, gamma)))

        u = np.random.uniform()
        if u < acceptance_prob:
            samples.append(proposed_sample)
            accepted_samples += 1
        else:
            samples.append(current_sample)

    acceptance_rate = accepted_samples / num_samples
    return np.array(samples), acceptance_rate

# Define the target density function f(x)
sigma = 0.1
delta = 1
target_density = lambda x: f(x, sigma, delta)

```

```

[6]: # Define the proposed density function g(y)
gamma_values = np.logspace(np.log10(0.01), np.log10(8), num=10)

print(gamma_values)
proposed_density = lambda y, gamma: g(y, gamma)

# Set the number of samples to generate
num_samples = 1000

```

```

[0.01      0.02101675 0.04417038 0.09283178 0.19510222 0.41004146
 0.86177388 1.81116859 3.80648772 8.          ]

```

```

[7]: # Set the initial sample
initial_sample = np.array([1,1]) # Adjust as needed

# Generate samples using Hastings-Metropolis sampling for different gamma values
samples = []
acceptance_rates = []
for gamma in gamma_values:
    result = hastings_metropolis_sampling(target_density, proposed_density,
    ↪initial_sample, num_samples, gamma)
    samples.append(result[0])

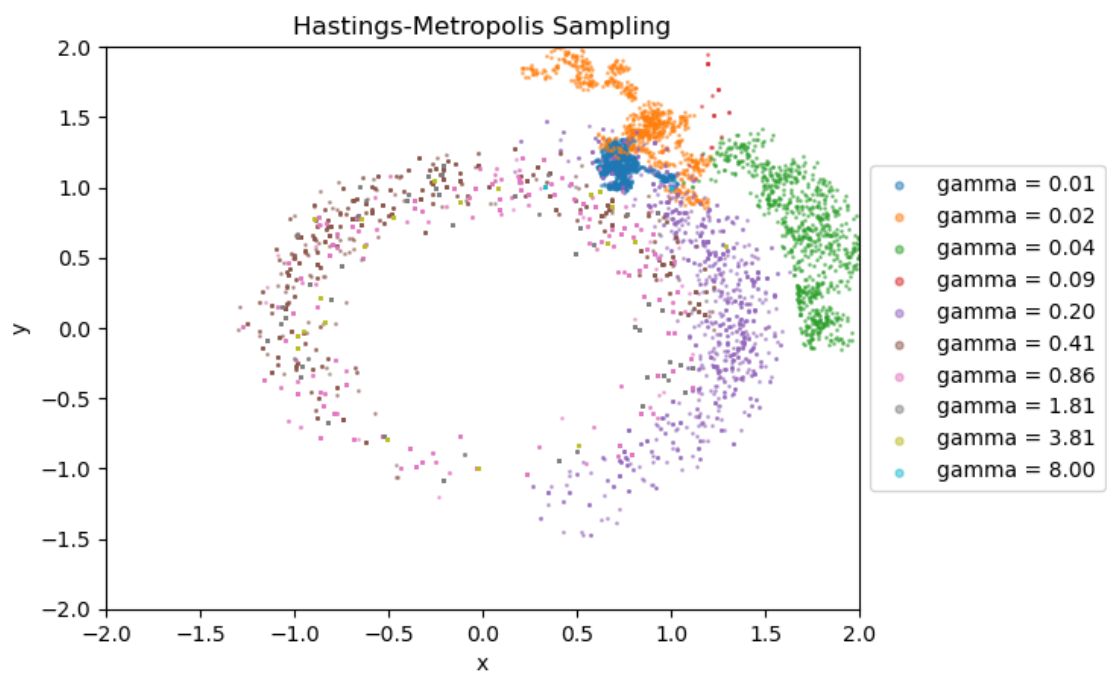
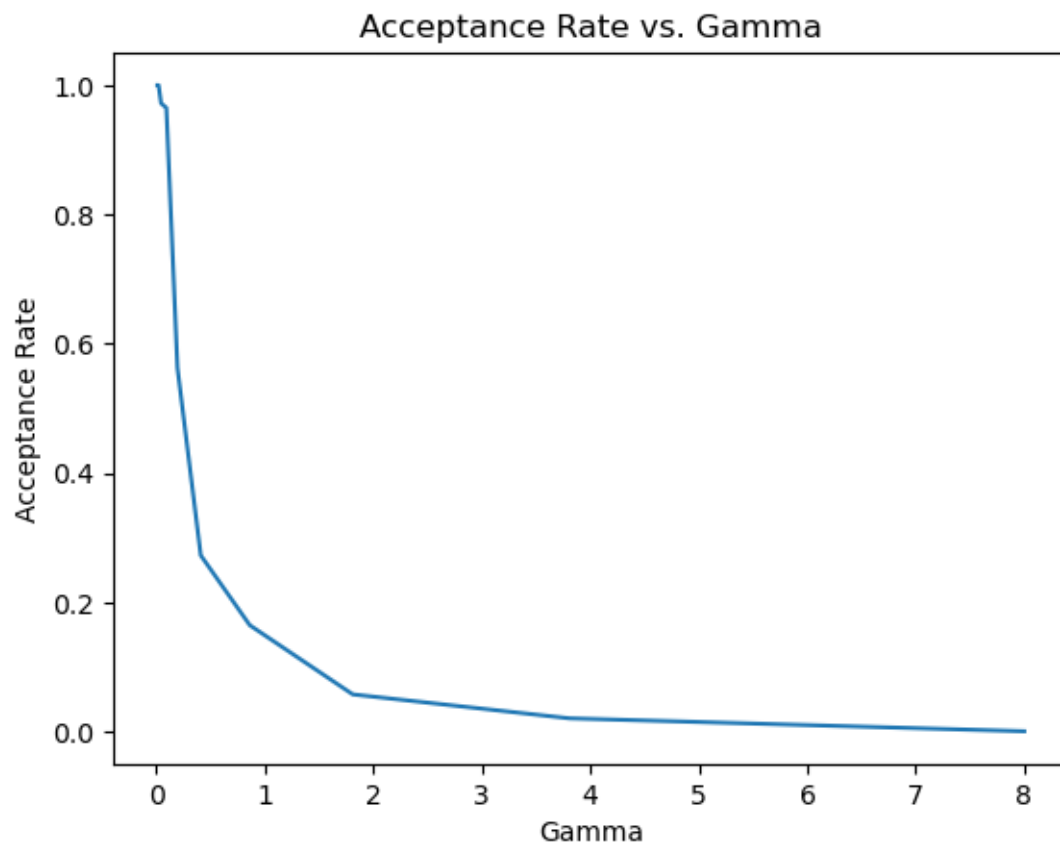
```

```
acceptance_rates.append(result[1])
```

```
/var/folders/22/nmmy2bvn43dby4g5jh1r3jzw0000gn/T/ipykernel_10619/123529884.py:13
: RuntimeWarning: invalid value encountered in double_scalars
    acceptance_prob = min(1, (target_density(proposed_sample) *
proposed_density(current_sample, gamma)) /
/var/folders/22/nmmy2bvn43dby4g5jh1r3jzw0000gn/T/ipykernel_10619/123529884.py:13
: RuntimeWarning: divide by zero encountered in double_scalars
    acceptance_prob = min(1, (target_density(proposed_sample) *
proposed_density(current_sample, gamma)) /
```

```
[8]: # Plotting the acceptance rates
plt.plot(gamma_values, acceptance_rates)
plt.xlabel('Gamma')
plt.ylabel('Acceptance Rate')
plt.title('Acceptance Rate vs. Gamma')
plt.show()


# Plotting the samples
for i, gamma in enumerate(gamma_values):
    plt.scatter(samples[i][:, 0], samples[i][:, 1], s=1, alpha=0.5,
        label=f'gamma = {gamma:.2f}')
plt.xlim(-2, 2)
plt.ylim(-2, 2)
plt.xlabel('x')
plt.ylabel('y')
plt.title('Hastings-Metropolis Sampling')
plt.legend(markerscale=3.5, loc='center left', bbox_to_anchor=(1, 0.5))
plt.show()
```



```
[9]: ## https://www.pymc.io/projects/docs/en/stable/learn/core\_notebooks/  
      ↪ pymc\_overview.html#pymc-overview  
  
import numpy as np  
import pymc as pm  
  
print(f"Running on PyMC v{pm.__version__}")
```

Running on PyMC v5.4.1

```
[10]: import arviz as az  
import pandas as pd  
import matplotlib.pyplot as plt  
  
RANDOM_SEED = 8927  
rng = np.random.default_rng(RANDOM_SEED)  
# True parameter values  
alpha, sigma = 1, 1  
beta = [1, 2.5]  
  
# Size of dataset  
size = 100  
  
# Predictor variable  
X1 = np.random.randn(size)  
X2 = np.random.randn(size) * 0.2  
  
# Simulate outcome variable  
Y = alpha + beta[0] * X1 + beta[1] * X2 + rng.normal(size=size) * sigma  
  
fig, axes = plt.subplots(1, 2, sharex=True, figsize=(10, 4))  
axes[0].scatter(X1, Y, alpha=0.6)  
axes[1].scatter(X2, Y, alpha=0.6)  
axes[0].set_ylabel("Y")  
axes[0].set_xlabel("X1")  
axes[1].set_xlabel("X2");
```



```
[11]: basic_model = pm.Model()

with basic_model:
    # Priors for unknown model parameters
    alpha = pm.Normal("alpha", mu=0, sigma=10)
    beta = pm.Normal("beta", mu=0, sigma=10, shape=2)
    sigma = pm.HalfNormal("sigma", sigma=1)

    # Expected value of outcome
    mu = alpha + beta[0] * X1 + beta[1] * X2

    # Likelihood (sampling distribution) of observations
    Y_obs = pm.Normal("Y_obs", mu=mu, sigma=sigma, observed=Y)
```

```
[12]: with basic_model:
    # draw 1000 posterior samples
    idata = pm.sample(1000)
```

Auto-assigning NUTS sampler...

Initializing NUTS using jitter+adapt_diag...

Multiprocess sampling (4 chains in 4 jobs)

NUTS: [alpha, beta, sigma]

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

Sampling 4 chains for 1_000 tune and 1_000 draw iterations (4_000 + 4_000 draws total) took 1 seconds.

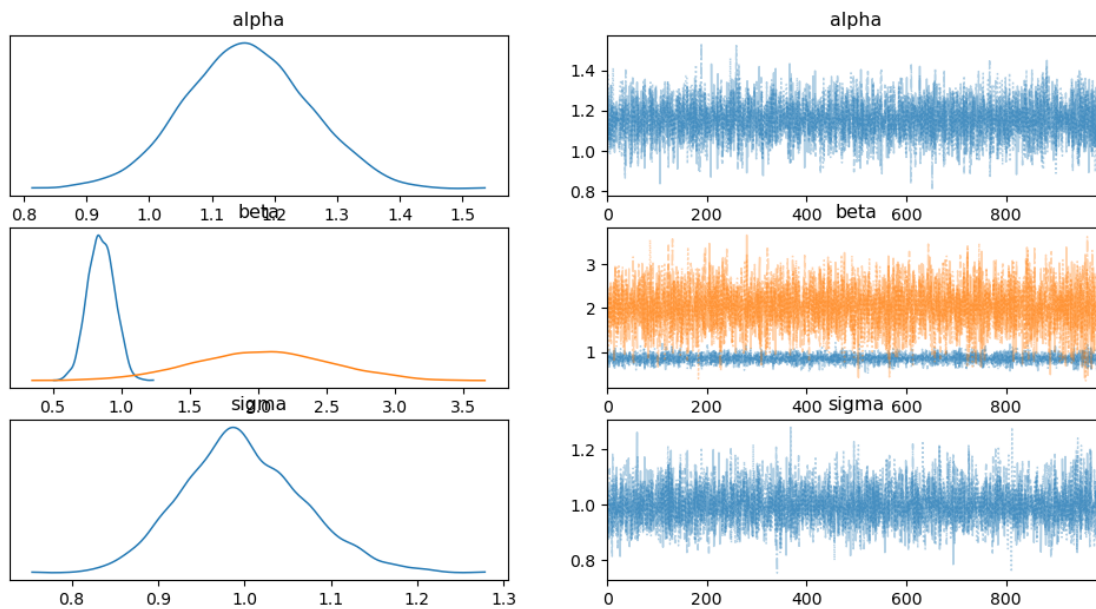
```
[13]: idata.posterior["alpha"].sel(draw=slice(0, 4))
```

```
[13]: <xarray.DataArray 'alpha' (chain: 4, draw: 5)>
array([[1.09940941, 1.31344334, 1.20623435, 1.08871712, 1.0559109 ],
       [1.23086922, 1.36363135, 1.37651772, 0.99086544, 1.27511779],
       [1.36839289, 1.28650864, 1.08145372, 1.11220016, 1.04912181],
       [1.02794591, 1.26310488, 1.05876462, 1.05876462, 1.15713395]])
```

Coordinates:

```
* chain    (chain) int64 0 1 2 3
* draw     (draw) int64 0 1 2 3 4
```

```
[14]: az.plot_trace(idata, combined=True);
```



```
[15]: az.summary(idata, round_to=2)
```

```
[15]:
```

	mean	sd	hdi_3%	hdi_97%	mcse_mean	mcse_sd	ess_bulk	ess_tail	\
alpha	1.16	0.10	0.98	1.34	0.00	0.0	6626.42	3198.90	
beta[0]	0.85	0.10	0.68	1.04	0.00	0.0	5746.07	3456.17	
beta[1]	2.03	0.50	1.09	2.97	0.01	0.0	5960.64	3143.41	
sigma	1.00	0.07	0.87	1.13	0.00	0.0	6039.31	3261.14	

	r_hat
alpha	1.0
beta[0]	1.0
beta[1]	1.0
sigma	1.0

```
[ ]:
```