

Opt

April 27, 2023

1 STAT 207: Optimization

Optimization considers the problem

$$\begin{array}{ll} \text{minimize} & f(x) \\ \text{subject to} & \text{constraints on } x \end{array}$$

- Maximization or minimization
- Exact or iterative solutions
- Why is optimization important in statistics?
 - Maximum likelihood estimation (MLE).
 - Maximum a posteriori (MAP) estimation in Bayesian framework.
 - Machine learning: minimize a loss + certain regularization.
 - ...
- Commonly used optimization methods:
 - Newton type algorithms
 - quasi-Newton algorithm
 - expectation-maximization (EM) algorithm
 - majorization-minimization (MM) algorithm
 - convex programming with emphasis in statistical applications

1.1 Basic results

Suppose $f(x)$ is differentiable on the open set U :

- differential $df(x)$
- gradient $\nabla f(x)$
- second differential (Hessian) $d^2 f(x) = \nabla^2 f(x)$

(Fermat) Suppose a differentiable function $f(x)$ has a local minimum at the point y of the open set U . Then $\nabla f(x)$ vanishes at y .

- Stationary point y : $\nabla f(y) = 0$

NAS Proposition 11.2.3 Suppose a twice continuously differentiable function $f(x)$ has a local minimum at the point y of the open set U . Then $d^2f(x)$ is positive semidefinite at y . Conversely, if y is a stationary point and $d^2f(y)$ is positive definite, then y is a local minimum.

- A function f is coercive if $\lim_{\|x\|_2 \rightarrow \infty} f(x) = \infty$.

Example $f(x) = \frac{1}{2}x^T Ax + b^T x + c$

Example Show that the sample mean and sample variance are the MLE of the theoretical mean and variance of a random sample y_1, y_2, \dots, y_n from a multivariate normal distribution.

1.2 Convexity

A function $f : R^n \rightarrow R$ is **convex** if

- $\text{dom} f$ is a convex set: $\alpha x + (1 - \alpha)y \in \text{dom} f$ for all $x, y \in \text{dom} f$ and any $\alpha \in (0, 1)$, and
- $f(\alpha x + (1 - \alpha)y) \leq \alpha f(x) + (1 - \alpha)f(y)$ for all $x, y \in \text{dom} f$ and any $\alpha \in (0, 1)$.

f is **strictly convex** if the inequality is strict for all $x \neq y$ and α .

Supporting hyperplane inequality A differentiable function f is convex if and only if

$$f(y) \geq f(x) + df(x)(y - x)$$

for all $x, y \in \text{dom} f$.

Second-order condition for convexity A twice differentiable function f is convex if and only if $\nabla^2 f(x)$ is PSD for all $x \in \text{dom} f$. It is strictly convex if and only if $\nabla^2 f(x)$ is PD for all $x \in \text{dom} f$.

1.2.1 Convexity and global optima

Suppose f is a convex function.

- Any stationary point y , i.e., $\nabla f(y) = 0$, is a global minimum. (Proof: By supporting hyperplane inequality, $f(x) \geq f(y) + \nabla f(y)^\top(x - y) = f(y)$ for all $x \in \text{dom} f$.)
- Any local minimum is a global minimum.
- The set of (global) minima is convex.
- If f is strictly convex, then the global minimum, if exists, is unique.

Example: Least squares estimate. $f(\beta) = \frac{1}{2}\|y - X\beta\|_2^2$ has Hessian $\nabla^2 f = X^\top X$ which is positive semidefinite. So f is convex and any stationary point (solution to the normal equation) is a global minimum. When X is rank deficient, the set of solutions is convex.

1.2.2 Jensen's inequality

If h is convex and W a random vector taking values in $\text{dom} f$, then

$$\mathbb{E}[h(W)] \geq h[\mathbb{E}(W)],$$

provided both expectations exist. For a strictly convex, equality holds if and only if $W = \mathbf{E}(W)$ almost surely.

Proof: Take $x = W$ and $y = \mathbf{E}(W)$ in the supporting hyperplane inequality.

1.2.3 Information inequality

Let f and g be two densities with respect to a common measure μ and $f, g > 0$ almost everywhere relative to μ . Then

$$\mathbb{E}_f[\log(f)] \geq \mathbb{E}_f[\log(g)],$$

with equality if and only if $f = g$ almost everywhere on \mathcal{X} .

Proof: Apply Jensen's inequality to the convex function $-\ln(t)$ and random variable $X = (f/g)(\omega)$ where $\omega \sim \mu$.

Important applications of information inequality: M-estimation, EM algorithm.

1.3 Optimization with Equality Constraints

- Suppose the objective function $f(x)$ to be minimized is continuously differentiable and defined on \mathbb{R}^n .
- The gradient direction $\nabla f(x) = df(x)^T$ is the direction of steepest ascent of $f(x)$ near the point x .
- The following linear approximation is often used

$$f(x + su) = f(x) + s df(x)u + o(s),$$

for a unit vector u and a scalar s .

Lagrange multipliers

The Lagrangian function

$$\mathcal{L}(x, w) = f(x) + \sum_{i=1}^m w_i g_i(x),$$

where $f(x)$ is the objective function to minimize, $g_i(x) = 0$ are equality constraints.




FIGURE 11.1. Level Curves and Steepest Ascent Directions for $\sin(x)\cos(y)$

Proposition 11.3.1 (Lagrange) Suppose the continuously differentiable function $f(x)$ has a local minimum at the feasible point y and that the constraint functions $g_1(x), \dots, g_m(x)$ are continuously differentiable with linearly independent gradient vectors $\nabla g_i(y)$ at y . Then

- there exists a multiplier vector λ such that (y, λ) is a stationary point of the Lagrangian.

Furthermore, if $f(x)$ and all $g_i(x)$ are twice continuously differentiable, then

- $v^T \nabla^2 L(y) v \geq 0$ for every tangent vector v at y .

Conversely, if (y, λ) is a stationary point of the Lagrangian and $v^T \nabla^2 L(y) v > 0$ for every nontrivial tangent vector v at y , then

- y represents a local minimum of $f(x)$ subject to the constraints.

Example Quadratic Programming with Equality Constraints

Minimizing a quadratic function

$$q(x) = \frac{1}{2}x^T A x + b^T x + c$$

on \mathbb{R}^n subject to the m linear equality constraints

$$v_i^T x = d_i, \quad i = 1, \dots, m$$

is one of the most important problems in nonlinear programming.

To minimize $q(x)$ subject to the constraints, we introduce the Lagrangian

$$L(x, \lambda) = \frac{1}{2}x^T A x + b^T x + c + \sum_{i=1}^m \lambda_i (v_i^T x - d_i) = \frac{1}{2}x^T A x + b^T x + c + \lambda^T (Vx - d).$$

A stationary point of $L(x, \lambda)$ is determined by the equations

$$Ax + b + V^T \lambda = 0, Vx = d,$$

whose formal solution amounts to

$$\begin{pmatrix} x \\ \lambda \end{pmatrix} = \begin{pmatrix} A & V^T \\ V & 0 \end{pmatrix}^{-1} \begin{pmatrix} -b \\ d \end{pmatrix}.$$

The inverse exists thanks to the following proposition.

Proposition 11.3.2 Let A be an $n \times n$ positive definite matrix and V be an $m \times n$ matrix. Then the matrix

$$M = \begin{bmatrix} A & V^\top \\ V & 0 \end{bmatrix}$$

has inverse

$$M^{-1} = \begin{bmatrix} A^{-1} - A^{-1}V^\top(VA^{-1}V^\top)^{-1}VA^{-1} & -(VA^{-1}V^\top)^{-1}VA^{-1} \\ -A^{-1}V^\top(VA^{-1}V^\top)^{-1} & (VA^{-1}V^\top)^{-1} \end{bmatrix}$$

if and only if V has linearly independent rows v_{t_1}, \dots, v_{t_m} .

Additional **Example 11.3.4**

1.4 Optimization with Inequality Constraints

Minimize an objective function $f_0(x)$ subject to the mixed constraints

$$h_i(x) = 0, \quad 1 \leq i \leq p, \quad f_j(x) \leq 0, \quad 1 \leq j \leq m.$$

A constraint $f_j(x)$ is

- active at the feasible point x provided $f_j(x) = 0$;
- it is inactive if $f_j(x) < 0$.

To avoid redundant constraints, we need

- linear independence of the gradients of the equality constraints
- and a restriction on the active inequality constraints.

1.4.1 Duality

- The Lagrangian is

$$\mathcal{L}(\mathbf{x}, \lambda, \nu) = f_0(\mathbf{x}) + \sum_{i=1}^m \lambda_i f_i(\mathbf{x}) + \sum_{i=1}^p \nu_i h_i(\mathbf{x}).$$

The vectors $\lambda = (\lambda_1, \dots, \lambda_m)^\top$ and $\nu = (\nu_1, \dots, \nu_p)^\top$ are called the Lagrange multiplier vectors or dual variables.

- The **Lagrange dual function** is the minimum value of the Lagrangian over \mathbf{x} :

$$g(\lambda, \nu) = \inf_{\mathbf{x}} \mathcal{L}(\mathbf{x}, \lambda, \nu) = \inf_{\mathbf{x}} \left(f_0(\mathbf{x}) + \sum_{i=1}^m \lambda_i f_i(\mathbf{x}) + \sum_{i=1}^p \nu_i h_i(\mathbf{x}) \right).$$

- Denote the optimal value of the original problem by p^* . For any $\lambda \succeq 0$ and any ν , we have

$$g(\lambda, \nu) \leq p^*.$$

Proof: For any feasible point $\tilde{\mathbf{x}}$,

$$\mathcal{L}(\tilde{\mathbf{x}}, \lambda, \nu) = f_0(\tilde{\mathbf{x}}) + \sum_{i=1}^m \lambda_i f_i(\tilde{\mathbf{x}}) + \sum_{i=1}^p \nu_i h_i(\tilde{\mathbf{x}}) \leq f_0(\tilde{\mathbf{x}})$$

because the second term is non-positive and the third term is zero. Then,

$$g(\lambda, \nu) = \inf_{\mathbf{x}} \mathcal{L}(\mathbf{x}, \lambda, \nu) \leq \mathcal{L}(\tilde{\mathbf{x}}, \lambda, \nu) \leq f_0(\tilde{\mathbf{x}}).$$

- Since each pair (λ, ν) with $\lambda \geq 0$ gives a lower bound to the optimal value p^* , it is natural to ask for the best possible lower bound the Lagrange dual function can provide. This leads to the **Lagrange dual problem**

$$\max_{\lambda \geq 0} g(\lambda, \nu),$$

which is a convex problem regardless of whether the primal problem is convex or not.

- We denote the optimal value of the Lagrange dual problem by d^* , which satisfies the weak duality

$$d^* \leq p^*.$$

The difference $p^* - d^*$ is called the optimal duality gap.

- If the primal problem is convex, that is

$$\min_{\mathbf{x}} f_0(\mathbf{x}) \text{ subject to } f_i(\mathbf{x}) \leq 0, \quad i = 1, \dots, m, \mathbf{Ax} = \mathbf{b},$$

with f_0, \dots, f_m convex, we usually (but not always) have strong duality, i.e., $d^* = p^*$.

- The conditions under which strong duality holds are called constraint qualifications. A commonly used one is Slater's condition: There exists a point in the relative interior of the domain such that

$$f_i(\mathbf{x}) < 0, \quad i = 1, \dots, m, \mathbf{Ax} = \mathbf{b}.$$

Such a point is also called **strictly feasible**.

1.5 KKT (Karush, Kuhn, and Tucker) Conditions

- “One of the great triumphs of 20th century applied mathematics.”
- Original paper: [Nonlinear Programming by Kuhn and Tucker 1951](#)

1.5.1 Nonconvex problems

- Assume $f_0, \dots, f_m, h_1, \dots, h_p$ are differentiable. Let \mathbf{x}^* and (λ^*, ν^*) be any primal and dual optimal points with zero duality gap, i.e., strong duality holds.
- Since \mathbf{x}^* minimizes $L(\mathbf{x}, \lambda^*, \nu^*)$ over \mathbf{x} , its gradient vanishes at \mathbf{x}^* , we have the Karush-Kuhn-Tucker (KKT) conditions:

– Stationarity

$$\nabla f_0(\mathbf{x}^*) + \sum_{i=1}^m \lambda_i^* \nabla f_i(\mathbf{x}^*) + \sum_{i=1}^p \nu_i^* \nabla h_i(\mathbf{x}^*) = \mathbf{0}.$$

– Primal feasibility

$$f_i(\mathbf{x}^*) \leq 0, \quad i = 1, \dots, m, h_i(\mathbf{x}^*) = 0, \quad i = 1, \dots, p$$

– Dual feasibility

$$\lambda^* \geq \mathbf{0},$$

– Complementary slackness

$$\lambda^* \cdot \mathbf{f}(\mathbf{x}^*) = 0$$

- The fourth condition (complementary slackness) follows from:

$$f_0(\mathbf{x}^*) = \inf_{\mathbf{x}} \{f_0(\mathbf{x}) + \sum_{i=1}^m \lambda_i^* f_i(\mathbf{x}) + \sum_{i=1}^p \nu_i^* h_i(\mathbf{x})\}$$

Since $\sum_{i=1}^m \lambda_i^* f_i(\mathbf{x}^*) = 0$ and each term is non-positive, we have $\lambda_i^* f_i(\mathbf{x}^*) = 0$, $i = 1, \dots, m$.

- To summarize, for any optimization problem with differentiable objective and constraint functions for which strong duality obtains, any pair of primal and dual optimal points must satisfy the KKT conditions.

1.5.2 Convex problems

- When the primal problem is convex, the KKT conditions are also sufficient for the points to be primal and dual optimal.
- If f_i are convex and h_i are affine, and $(\tilde{x}, \tilde{\lambda}, \tilde{\nu})$ satisfy the KKT conditions, then \tilde{x} and $(\tilde{\lambda}, \tilde{\nu})$ are primal and dual optimal, with zero duality gap.
- The KKT conditions play an important role in optimization. Many algorithms for convex optimization are conceived as, or can be interpreted as, methods for solving the KKT conditions.